

A dark blue vertical bar on the left side of the page. A blue arrow points to the right from the bar, containing the year 2024.

2024

# Terraform

Architecting IAC Solutions for AWS

Several thin, curved lines in dark blue and light gray originate from the bottom left corner and sweep upwards and to the right, creating a sense of motion or growth.

# Contents

Introduction to Infrastructure as Code (IaC):.....	3
Understanding the benefits of IaC .....	3
Comparing IaC tools (Terraform, CloudFormation, Ansible, etc.) .....	3
Introduction to Terraform: .....	3
Core concepts: modules, providers, states, and workspaces .....	3
Installation and configuration of Terraform .....	4
Writing basic Terraform configuration files.....	4
AWS Fundamentals: .....	6
Overview of AWS services.....	6
Key concepts: regions, availability zones, EC2 instances, S3 buckets, VPCs, security groups, IAM roles and policies.....	6
Hands-on Lab: Creating a Simple EC2 Instance.....	6
Set Up the Project Directory .....	6
Write the Terraform Configuration File .....	7
Initialize the Terraform Working Directory.....	8
Validate the Configuration.....	8
Plan the Infrastructure Changes .....	8
Apply the Configuration.....	8
Verify the EC2 Instance .....	8
Clean Up Resources .....	8
Deep Dive into Terraform: .....	9
Variables and data sources .....	9
Modules and their organization.....	11
State management and remote backends.....	13
Terraform providers and their usage.....	15
AWS Networking with Terraform: .....	17
Creating VPCs, subnets, and internet gateways .....	17
Configuring route tables and network ACLs .....	18
Assigning public and private IP addresses to EC2 instances.....	19
Hands-on Lab: Creating a VPC and Subnets.....	20
Terraform Best Practices:.....	23
Modularizing Terraform configurations.....	23
Testing and validation of Terraform code .....	24
Security considerations in Terraform .....	25
Version control for Terraform code .....	26

Advanced AWS Services with Terraform: .....	27
Deploying S3 buckets and object storage .....	27
Creating IAM roles and policies .....	28
Configuring security groups and network ACLs .....	29
Automating deployments with Terraform Cloud.....	31
Hands-on Lab: Creating an S3 Bucket with Versioning .....	31
Write a Terraform Configuration File to Create an S3 Bucket with Versioning.....	31
Apply the Configuration to Create the S3 Bucket.....	32
Terraform and AWS Security: .....	33
Implementing security best practices in Terraform configurations .....	33
Securing AWS resources with IAM roles and policies.....	34
Using security groups to control network traffic.....	35
Real-world Use Cases: .....	37
Deploying multi-tier applications on AWS.....	37
Automating infrastructure for CI/CD pipelines.....	39
Implementing infrastructure as code for disaster recovery .....	41
Hands-on Lab: Creating a Multi-Tier Application.....	42

# Introduction to Infrastructure as Code (IaC):

## Understanding the benefits of IaC

1. **Consistency:** Use declarative scripts to ensure identical infrastructure across environments.  
*Example: Deploy identical environments for development, staging, and production.*
2. **Automation:** Reduces manual configuration tasks.  
*Example: Automate the provisioning of a new server with predefined specs.*
3. **Version Control:** Infrastructure can be tracked and rolled back via Git.  
*Example: Use GitHub to track updates to your infrastructure scripts.*
4. **Scalability:** Easily replicate or scale infrastructure.  
*Example: Scale a load balancer configuration to handle increased traffic.*
5. **Cost Management:** Better visibility into provisioned resources.  
*Example: Monitor and destroy unused resources to save costs.*

## Comparing IaC tools (Terraform, CloudFormation, Ansible, etc.)

### Terraform

- **Strength:** Cloud-agnostic.
- **Example:** Use Terraform to provision EC2 instances on AWS and Compute Engine on GCP.

### AWS CloudFormation

- **Strength:** AWS-specific.
- **Example:** Define an S3 bucket, an EC2 instance, and a VPC in YAML/JSON.

### Ansible

- **Strength:** Configuration management.
- **Example:** Install software on an already-provisioned EC2 instance.

### Pulumi

- **Strength:** Uses familiar programming languages.
- **Example:** Write Python code to create an AWS Lambda function.

# Introduction to Terraform:

## Core concepts: modules, providers, states, and workspaces

- a) **Modules:** Reusable, self-contained infrastructure components.  
*Example: A module for creating a VPC with subnets and routing tables.*
- b) **Providers:** Plugins that interact with cloud providers.  
*Example: provider "aws" { region = "us-east-1" } to configure AWS as a provider.*
- c) **State Files:** Keep track of infrastructure and its current state.  
*Example: A terraform.tfstate file showing deployed EC2 instances.*

- d) Workspaces: Isolate state files for different environments.  
*Example: Use default for staging and prod workspace for production.*

## Installation and configuration of Terraform

1. Download Terraform from the official site.
2. Add it to the system PATH.
3. Verify the installation:

```
terraform --version
```

## Writing basic Terraform configuration files

Terraform uses a declarative language to define infrastructure as code. A Terraform configuration file typically consists of provider configuration, resource definitions, and optional variables, outputs, and modules.

### Step 1: Set Up Terraform Environment

Install Terraform:

- a) Download Terraform from the Terraform website.
- b) Add Terraform to your system's PATH.

Create a Directory:

- a) Create a folder to store your Terraform configuration files. For example:

```
mkdir terraform-basic-config  
cd terraform-basic-config
```

### Step 2: Create the main.tf File

- a) Create a file named main.tf in your project directory. This file will contain the configuration.

```
# Provider Configuration  
provider "aws" {  
  region = "us-east-1"  
}  
  
# Resource Configuration: Create an EC2 Instance  
resource "aws_instance" "example" {  
  ami           = "ami-0c55b159cbf1f0" # Replace with a valid AMI ID  
  instance_type = "t2.micro"  
  
  tags = {  
    Name = "BasicTerraformInstance"  
  }  
}
```

### Step 3: Initialize Terraform

- a) Run the following command to download the necessary provider plugins and prepare the directory:

```
terraform init
```

\*This initializes Terraform and downloads the AWS provider plugin.

#### Step 4: Validate the Configuration

- a) Check if your configuration is valid by running:

```
terraform validate
```

Output:

Success! The configuration is valid.

#### Step 5: Plan the Infrastructure

- a) Preview the changes Terraform will make:

```
terraform plan
```

Output includes the resources Terraform will create (e.g., EC2 instance).

#### Step 6: Apply the Configuration

- a) Deploy the defined infrastructure:

```
terraform apply
```

\*Type yes when prompted.

Output:

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

#### Step 7: Verify the Deployed Resource

- a) Log in to your AWS Management Console.  
b) Navigate to the EC2 section in the specified region (us-east-1) to confirm the instance has been created.

#### Step 8: Clean Up

- a) Destroy the resources when you're done:

```
terraform destroy
```

Type yes to confirm.

### Additional Examples

#### Example 1: Create an S3 Bucket

```
resource "aws_s3_bucket" "example" {  
  bucket = "my-unique-terraform-bucket"  
  acl    = "private"  
}
```

#### Example 2: Use Variables

```
variable "region" {  
  default = "us-east-1"  
}  
  
provider "aws" {  
  region = var.region  
}
```

#### Example 3: Add an EBS Volume

```
resource "aws_ebs_volume" "example" {
  availability_zone = "us-east-1a"
  size             = 10
}
```

#### Example 4: Use Output Values

```
output "instance_id" {
  value = aws_instance.example.id
}
```

#### Example 5: Use a Data Source

```
data "aws_ami" "latest" {
  most_recent = true
  owners      = ["amazon"]

  filter {
    name   = "name"
    values = ["amzn2-ami-hvm-*-x86_64-gp2"]
  }
}

resource "aws_instance" "example" {
  ami          = data.aws_ami.latest.id
  instance_type = "t2.micro"
}
```

## AWS Fundamentals:

### Overview of AWS services

- *See attached AWS Fundamentals Trainer's Guide for this topic.*
- *Topic will be discussed by trainer separately*

Key concepts: regions, availability zones, EC2 instances, S3 buckets, VPCs, security groups, IAM roles and policies

- *See attached AWS Fundamentals Trainer's Guide for this topic.*
- *Topic will be discussed by trainer separately*

## Hands-on Lab: Creating a Simple EC2 Instance

### Set Up the Project Directory

Create a new folder for your Terraform project:

```
mkdir terraform-ec2-lab
cd terraform-ec2-lab
```

Inside this folder, create a new file named main.tf:

```
touch main.tf
```

Write the Terraform Configuration File

Edit the main.tf file and add the following content:

**main.tf Content:**

```
# Configure the AWS Provider
provider "aws" {
  region = "us-east-1" # Replace with your desired AWS region
}

# Create a Key Pair
resource "aws_key_pair" "terraform_key" {
  key_name   = "terraform-key"
  public_key = file("~/ssh/id_rsa.pub") # Use your local SSH public key file
}

# Create a Security Group
resource "aws_security_group" "allow_ssh" {
  name_prefix = "allow-ssh-"

  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

# Create an EC2 Instance
resource "aws_instance" "example" {
  ami          = "ami-0c55b159cbf0" # Replace with a valid AMI ID for your region
  instance_type = "t2.micro"

  key_name       = aws_key_pair.terraform_key.key_name
  security_groups = [aws_security_group.allow_ssh.name]

  tags = {
    Name = "Terraform-EC2-Instance"
  }
}
```



## Initialize the Terraform Working Directory

Run the following command to initialize Terraform and download the necessary provider plugins:

```
terraform init
```

Expected Output:

Terraform has been successfully initialized!

## Validate the Configuration

Validate the configuration to ensure there are no syntax errors:

```
terraform validate
```

Expected Output:

Success! The configuration is valid.

## Plan the Infrastructure Changes

Run a terraform plan to preview the infrastructure changes Terraform will make:

```
terraform plan
```

Expected Output:

A list of resources to be created, such as the EC2 instance, security group, and key pair.

## Apply the Configuration

Deploy the infrastructure:

```
terraform apply
```

- ✓ Terraform will show you a plan similar to the one generated by terraform plan.
- ✓ Type yes to confirm and create the resources.

Expected Output:

Apply complete! Resources: 3 added, 0 changed, 0 destroyed.

## Verify the EC2 Instance

- 1) Log in to the AWS Management Console.
- 2) Navigate to EC2 > Instances in the us-east-1 region (or the region you specified).
- 3) Confirm that the EC2 instance is running, and check its public IP address.

## Clean Up Resources

When you're done, destroy the infrastructure to avoid incurring costs:

```
terraform destroy
```

\*Type yes to confirm resource deletion.

Verify the output:

Destroy complete! Resources: 3 destroyed.

## Key Features Demonstrated:

Provider Configuration: Set up AWS as the provider.

Resource Management:

- Created a key pair for SSH access.
- Configured a security group to allow SSH traffic.
- Provisioned an EC2 instance using a specific AML.

Lifecycle Commands:

terraform init	-to initialize the environment.
terraform plan	-to preview changes.
terraform apply	-to deploy the infrastructure.
terraform destroy	-to clean up resources.

## Deep Dive into Terraform:

Variables and data sources

### Variables

Variables allow dynamic values in your Terraform configuration.

Define a Variable:

In a file called variables.tf:

<pre>variable "instance_type" {   default = "t2.micro"   description = "The type of EC2 instance to create" }</pre>
---

Use the Variable:

In your main.tf file:

<pre>resource "aws_instance" "example" {   ami      = "ami-0c55b159cbf0e1f0"   instance_type = var.instance_type }</pre>
--

Set the Variable Value:

<pre>terraform apply -var="instance_type=t3.medium"</pre>
---

Or use a terraform.tfvars file:

<pre>instance_type = "t3.medium"</pre>
--

## Data Sources

Data sources fetch information about existing resources.

### Define a Data Source:

```
data "aws_ami" "latest" {
  most_recent = true
  owners     = ["amazon"]

  filter {
    name = "name"
    values = ["amzn2-ami-hvm-*-x86_64-gp2"]
  }
}
```

### Use the Data Source:

```
resource "aws_instance" "example" {
  ami      = data.aws_ami.latest.id
  instance_type = "t2.micro"
}
```

## Example 1: Dynamic EC2 Instance Type

This example dynamically sets the instance type using variables and fetches the latest Amazon Linux AMI using a data source.

### Configuration:

```
# variables.tf
variable "instance_type" {
  default = "t2.micro"
  description = "The type of EC2 instance"
}
```

```
# data.tf
data "aws_ami" "latest" {
  most_recent = true
  owners     = ["amazon"]

  filter {
    name = "name"
    values = ["amzn2-ami-hvm-*-x86_64-gp2"]
  }
}
```

```
# main.tf
resource "aws_instance" "example" {
  ami      = data.aws_ami.latest.id
  instance_type = var.instance_type

  tags = {
    Name = "DynamicEC2Instance"
  }
}
```

```
}  
}
```

Steps:

- Change the instance\_type dynamically via CLI or terraform.tfvars.
- The AMI is fetched dynamically using the data source.

### Example 2: Dynamic S3 Bucket Name

This example creates an S3 bucket with a name derived from a variable.

Configuration:

```
# variables.tf  
variable "project_name" {  
  description = "Project name for the bucket"  
  default    = "my-project"  
}
```

```
# main.tf  
resource "aws_s3_bucket" "example" {  
  bucket = "${var.project_name}-bucket"  
  acl    = "private"  
}  
  
data "aws_region" "current" {}  
  
output "bucket_region" {  
  value = data.aws_region.current.name  
}
```

Steps:

- Change the project\_name variable to modify the bucket name.
- Use the aws\_region data source to fetch the current region for display.

### Modules and their organization

Modules help organize reusable Terraform code.

#### Create a Module

Create a directory for your module:

```
mkdir -p modules/ec2  
cd modules/ec2
```

Add main.tf:

```
resource "aws_instance" "example" {  
  ami      = var.ami  
  instance_type = var.instance_type  
  
  tags = {  
    Name = var.instance_name  
  }
```

```
}  
}
```

Add variables.tf:

```
variable "ami" {}  
variable "instance_type" {  
  default = "t2.micro"  
}  
variable "instance_name" {}
```

Add outputs.tf:

```
output "instance_id" {  
  value = aws_instance.example.id  
}
```

### Use the Module

In your main project directory, reference the module:

```
module "ec2_instance" {  
  source      = "./modules/ec2"  
  ami         = "ami-0c55b159cbfafa1f0"  
  instance_type = "t2.micro"  
  instance_name = "MyEC2Instance"  
}
```

Run Terraform:

```
terraform init  
terraform apply
```

### **Example 1: VPC Module**

This example uses a module to create a VPC with subnets and a route table.

Module (in modules/vpc/main.tf):

```
resource "aws_vpc" "example" {  
  cidr_block = var.cidr_block  
}  
  
resource "aws_subnet" "example" {  
  count      = var.subnet_count  
  vpc_id    = aws_vpc.example.id  
  cidr_block = cidrsubnet(var.cidr_block, 8, count.index)  
}  
  
variable "cidr_block" {}  
variable "subnet_count" {}
```

Main Configuration:

```
module "vpc" {  
  source    = "../modules/vpc"  
  cidr_block = "10.0.0.0/16"  
  subnet_count = 2  
}
```

### Example 2: EC2 Module

Reuse an EC2 module for different environments.

Module (in modules/ec2/main.tf):

```
resource "aws_instance" "example" {  
  ami          = var.ami  
  instance_type = var.instance_type  
  
  tags = {  
    Name = var.instance_name  
  }  
}  
  
variable "ami" {}  
variable "instance_type" {}  
variable "instance_name" {}
```

Main Configuration:

```
module "staging_ec2" {  
  source    = "../modules/ec2"  
  ami       = "ami-0c55b159cbfafa1f0"  
  instance_type = "t2.micro"  
  instance_name = "StagingInstance"  
}  
  
module "prod_ec2" {  
  source    = "../modules/ec2"  
  ami       = "ami-0c55b159cbfafa1f0"  
  instance_type = "t3.medium"  
  instance_name = "ProdInstance"  
}
```

State management and remote backends

#### Understand State

Terraform tracks infrastructure in a .tfstate file. By default, this file is stored locally.

#### Enable a Remote Backend

Configure backend in your main.tf:

```
terraform {  
  backend "s3" {  
    bucket = "my-terraform-state"  
    key    = "state/terraform.tfstate"  
  }  
}
```

```
    region    = "us-east-1"
  }
}
```

Initialize Terraform:

```
terraform init
```

\*This migrates the local state to the remote backend.

### Example 1: Remote Backend in S3

This example sets up Terraform state storage in an S3 bucket.

Main Configuration:

```
terraform {
  backend "s3" {
    bucket    = "my-terraform-state"
    key       = "state/terraform.tfstate"
    region    = "us-east-1"
  }
}

provider "aws" {
  region = "us-east-1"
}
```

Steps:

- Create the S3 bucket my-terraform-state.
- Run terraform init to migrate the local state to the S3 backend.

### Example 2: Shared State for Multiple Teams

Use a remote backend and locking for collaboration.

Main Configuration:

```
terraform {
  backend "s3" {
    bucket    = "shared-terraform-state"
    key       = "dev/terraform.tfstate"
    region    = "us-east-1"
    dynamodb_table = "terraform-locking"
  }
}
```

Steps:

- Configure DynamoDB table terraform-locking for state locking.
- Teams working on the same configuration will avoid race conditions.

## Terraform providers and their usage

### What Are Providers?

Providers are plugins used to interact with cloud platforms and other APIs.

### Configure a Provider

Add a provider block:

```
provider "aws" {  
  region = "us-east-1"  
}
```

Specify provider requirements in versions.tf:

```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = "~> 4.0"  
    }  
  }  
}
```

### Use Multiple Providers

Define multiple providers:

```
provider "aws" {  
  alias = "primary"  
  region = "us-east-1"  
}  
  
provider "aws" {  
  alias = "secondary"  
  region = "us-west-2"  
}
```

Reference a provider in resources:

```
resource "aws_instance" "example" {  
  provider    = aws.primary  
  ami        = "ami-0c55b159cbfafa1f0"  
  instance_type = "t2.micro"  
}
```

### **Example 1: AWS and GCP Providers**

Provision an EC2 instance in AWS and a VM in Google Cloud.

Configuration:

```
provider "aws" {  
  region = "us-east-1"  
}  
  
provider "google" {  
  project = "my-gcp-project"}
```



```

    region = "us-central1"
  }

resource "aws_instance" "example" {
  ami      = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
}

resource "google_compute_instance" "example" {
  name      = "gcp-vm"
  machine_type = "n1-standard-1"
  zone      = "us-central1-a"

  boot_disk {
    initialize_params {
      image = "debian-cloud/debian-11"
    }
  }

  network_interface {
    network = "default"
  }
}

```

### Example 2: Multiple AWS Regions

Deploy instances in different AWS regions using provider aliases.

Configuration:

```

provider "aws" {
  alias = "us_east"
  region = "us-east-1"
}

provider "aws" {
  alias = "us_west"
  region = "us-west-2"
}

resource "aws_instance" "east_instance" {
  provider = aws.us_east
  ami      = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
}

resource "aws_instance" "west_instance" {
  provider = aws.us_west
  ami      = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
}

```

## AWS Networking with Terraform:

### Creating VPCs, subnets, and internet gateways

A VPC (Virtual Private Cloud) is a logically isolated network for AWS resources. Subnets divide the VPC into smaller segments, and Internet Gateways provide internet connectivity.

#### Example 1: Creating a Simple VPC with Subnets

Configuration:

```
resource "aws_vpc" "example" {
  cidr_block      = "10.0.0.0/16"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = {
    Name = "MyVPC"
  }
}

resource "aws_subnet" "example" {
  vpc_id      = aws_vpc.example.id
  cidr_block   = "10.0.1.0/24"
  availability_zone = "us-east-1a"
  tags = {
    Name = "MySubnet"
  }
}
```

Steps:

- Run terraform apply.
- Check the VPC and subnet in the AWS console.

#### Example 2: VPC with Two Subnets

Configuration:

```
resource "aws_vpc" "example" {
  cidr_block = "10.0.0.0/16"
  tags = {
    Name = "MyVPC"
  }
}

resource "aws_subnet" "public_subnet" {
  vpc_id      = aws_vpc.example.id
  cidr_block   = "10.0.1.0/24"
  availability_zone = "us-east-1a"
  tags = {
    Name = "PublicSubnet"
  }
}

resource "aws_subnet" "private_subnet" {
```

```
vpc_id      = aws_vpc.example.id
cidr_block  = "10.0.2.0/24"
availability_zone = "us-east-1a"
tags = {
  Name = "PrivateSubnet"
}
```

### Example 3: Adding an Internet Gateway

Configuration:

```
resource "aws_vpc" "example" {
  cidr_block = "10.0.0.0/16"
  tags = {
    Name = "MyVPC"
  }
}

resource "aws_internet_gateway" "example" {
  vpc_id = aws_vpc.example.id
  tags = {
    Name = "MyInternetGateway"
  }
}

resource "aws_subnet" "example" {
  vpc_id      = aws_vpc.example.id
  cidr_block  = "10.0.1.0/24"
  availability_zone = "us-east-1a"
  tags = {
    Name = "MySubnet"
  }
}
```

Configuring route tables and network ACLs

### Example 1: Adding a Route Table for Internet Access

Configuration:

```
resource "aws_route_table" "example" {
  vpc_id = aws_vpc.example.id

  route {
    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.example.id
  }

  tags = {
    Name = "MyRouteTable"
  }
}
```

```
resource "aws_route_table_association" "example" {
  subnet_id    = aws_subnet.example.id
  route_table_id = aws_route_table.example.id
}
```

### Example 2: Creating a Network ACL

Configuration:

```
resource "aws_network_acl" "example" {
  vpc_id = aws_vpc.example.id

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    rule_no   = 100
    cidr_block = "0.0.0.0/0"
    action    = "allow"
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    rule_no   = 200
    cidr_block = "0.0.0.0/0"
    action    = "allow"
  }

  tags = {
    Name = "MyNetworkACL"
  }
}
```

### Example 3: Associating a Subnet with a Network ACL

Configuration:

```
resource "aws_network_acl_association" "example" {
  subnet_id    = aws_subnet.example.id
  network_acl_id = aws_network_acl.example.id
}
```

Assigning public and private IP addresses to EC2 instances

### Example 1: Assigning a Public IP Address

Configuration:

```
resource "aws_instance" "example" {
  ami          = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
  subnet_id    = aws_subnet.example.id
  associate_public_ip_address = true
}
```

```
tags = {
  Name = "PublicInstance"
}
```

### Example 2: Assigning a Private IP Address

Configuration:

```
resource "aws_instance" "example" {
  ami          = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
  subnet_id    = aws_subnet.private_subnet.id
  private_ip   = "10.0.2.10"

  tags = {
    Name = "PrivateInstance"
  }
}
```

### Example 3: Creating an EC2 Instance in Both Public and Private Subnets

Configuration:

```
module "public_instance" {
  source          = "./modules/ec2"
  subnet_id       = aws_subnet.public_subnet.id
  associate_public_ip_address = true
  instance_name   = "PublicInstance"
}

module "private_instance" {
  source          = "./modules/ec2"
  subnet_id       = aws_subnet.private_subnet.id
  private_ip      = "10.0.2.10"
  instance_name   = "PrivateInstance"
}
```

## Hands-on Lab: Creating a VPC and Subnets

### Set Up Your Project Directory

Create a new directory for your Terraform project:

```
mkdir terraform-vpc
cd terraform-vpc
```

Create the main Terraform configuration file:

```
touch main.tf
```

## Write the Terraform Configuration File

Here's how to create a VPC with one public and one private subnet.

Complete Configuration (main.tf):

```
# Provider Configuration
provider "aws" {
  region = "us-east-1" # Specify your AWS region
}

# Create a VPC
resource "aws_vpc" "my_vpc" {
  cidr_block      = "10.0.0.0/16"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = {
    Name = "MyVPC"
  }
}

# Create an Internet Gateway
resource "aws_internet_gateway" "my_igw" {
  vpc_id = aws_vpc.my_vpc.id
  tags = {
    Name = "MyInternetGateway"
  }
}

# Create a Public Subnet
resource "aws_subnet" "public_subnet" {
  vpc_id      = aws_vpc.my_vpc.id
  cidr_block   = "10.0.1.0/24"
  map_public_ip_on_launch = true
  availability_zone = "us-east-1a"
  tags = {
    Name = "PublicSubnet"
  }
}

# Create a Private Subnet
resource "aws_subnet" "private_subnet" {
  vpc_id      = aws_vpc.my_vpc.id
  cidr_block   = "10.0.2.0/24"
  availability_zone = "us-east-1a"
  tags = {
    Name = "PrivateSubnet"
  }
}

# Create a Route Table for Public Subnet
resource "aws_route_table" "public_route_table" {
```

```
vpc_id = aws_vpc.my_vpc.id
route {
  cidr_block = "0.0.0.0/0"
  gateway_id = aws_internet_gateway.my_igw.id
}
tags = {
  Name = "PublicRouteTable"
}
}
```

# Associate Public Subnet with the Route Table

```
resource "aws_route_table_association" "public_route_assoc" {
  subnet_id      = aws_subnet.public_subnet.id
  route_table_id = aws_route_table.public_route_table.id
}
```

### Initialize Terraform

Run the following command to initialize Terraform. This downloads necessary provider plugins:

```
terraform init
```

Expected Output:

Terraform has been successfully initialized!

### Plan the Infrastructure

Run terraform plan to preview the resources Terraform will create:

```
terraform plan
```

Expected Output:

A summary of the resources to be created, including the VPC, subnets, and internet gateway.

### Apply the Configuration

Run the following command to create the VPC and subnets:

```
terraform apply
```

\*When prompted, type yes to confirm.

Expected Output:

Apply complete! Resources: 5 added, 0 changed, 0 destroyed.

### Verify the Resources in AWS

- a) Log in to the AWS Management Console.
- b) Navigate to the VPC section.
- c) Verify that:
  - ✓ A VPC with CIDR block 10.0.0.0/16 exists.
  - ✓ One public subnet (10.0.1.0/24) is associated with the internet gateway.
  - ✓ One private subnet (10.0.2.0/24) is created.

## Clean Up Resources

To avoid unnecessary charges, destroy the resources when you're done:

```
terraform destroy
```

\*Type yes to confirm.

## Terraform Best Practices:

Modularizing Terraform configurations

Modularizing Terraform code involves breaking down the configuration into reusable, organized units called modules.

### Create a Module

Modules are self-contained configurations.

Directory Structure:

```
terraform-project/
├── main.tf      # Main file to call modules
├── variables.tf # Variables for the root module
├── outputs.tf   # Outputs for the root module
└── modules/
    └── vpc/
        ├── main.tf
        ├── variables.tf
        └── outputs.tf
```

Module (modules/vpc):

```
# modules/vpc/main.tf:
resource "aws_vpc" "example" {
  cidr_block = var.cidr_block
  tags = {
    Name = var.name
  }
}

resource "aws_subnet" "public_subnet" {
  vpc_id      = aws_vpc.example.id
  cidr_block  = var.public_subnet_cidr
  map_public_ip_on_launch = true
}
```

```
# modules/vpc/variables.tf:
variable "cidr_block" {}
variable "name" {}
variable "public_subnet_cidr" {}
```

```
# modules/vpc/outputs.tf:
output "vpc_id" {
  value = aws_vpc.example.id
}
```



```
output "public_subnet_id" {
  value = aws_subnet.public_subnet.id
}
```

Root Module (main.tf):

```
module "vpc" {
  source      = "./modules/vpc"
  cidr_block  = "10.0.0.0/16"
  name       = "MyVPC"
  public_subnet_cidr = "10.0.1.0/24"
}
```

## Apply Modular Configuration

Initialize Terraform:

```
terraform init
```

Plan and apply:

```
terraform plan
terraform apply
```

Benefits of Modularization:

- ✓ Reusability: Modules can be reused across projects.
- ✓ Maintainability: Easier to manage smaller pieces of code.
- ✓ Team Collaboration: Different teams can work on separate modules.

Testing and validation of Terraform code

Testing and validating Terraform ensures infrastructure is deployed as expected and reduces risks.

## Syntax Validation

Validate configuration files for syntax errors.

```
terraform validate
```

## Pre-Deployment Testing

Plan Command: Preview the infrastructure changes.

```
terraform plan
```

## Unit Testing with terratest

Use the terratest library for automated testing.

- 1) Install terratest in Go.
- 2) Write a test:

```
go
package test

import (
    "testing"
    "github.com/gruntwork-io/terratest/modules/terraform"
)
```

```
func TestTerraformVPC(t *testing.T) {
    options := &terraform.Options{
        TerraformDir: "../terraform-project",
    }

    defer terraform.Destroy(t, options)

    terraform.InitAndApply(t, options)

    vpcID := terraform.Output(t, options, "vpc_id")
    if vpcID == "" {
        t.Fatal("VPC ID is empty")
    }
}
```

Run the test:

```
go test -v
```

### Linting with tflint

Install tflint and run it to check for potential errors.

```
tflint
```

## Security considerations in Terraform

### Secure State Files

Store State Remotely: Use secure backends like S3 with encryption:

```
terraform {
    backend "s3" {
        bucket    = "my-terraform-state"
        key       = "state/terraform.tfstate"
        region    = "us-east-1"
        encrypt   = true
        dynamodb_table = "terraform-locking"
    }
}
```

\*Encrypt State Locally: Use secure storage for local .tfstate files (e.g., disk encryption).

### Manage Secrets

Avoid Hardcoding Secrets: Use environment variables or secret management tools like AWS Secrets Manager or HashiCorp Vault.

Example with Vault:

```
export VAULT_ADDR='http://127.0.0.1:8200'
export VAULT_TOKEN='my-token'
terraform plan
```

## Least Privilege IAM Roles

Assign minimum permissions to the Terraform IAM user or role.

## Use Security Scanners

Scan Terraform configurations with tools like:

```
checkov
tfsec
```

Version control for Terraform code

## Use Git for Version Control

Initialize a Git repository:

```
git init
```

Track Terraform files:

```
git add main.tf variables.tf outputs.tf
git commit -m "Initial Terraform configuration"
```

## Use Branching for Changes

Create a feature branch:

```
git checkout -b feature/add-vpc
```

Commit and push changes:

```
git commit -am "Add VPC configuration"
git push origin feature/add-vpc
```

## Use .gitignore

Ignore sensitive files and local state files:

```
*.tfstate
*.tfstate.backup
.terraform/
```

## Tagging for Releases

Tag stable versions of infrastructure:

```
git tag -a v1.0 -m "Stable release"
git push origin v1.0
```

## Automate with CI/CD

Use GitHub Actions or Jenkins to automate Terraform workflows.

Example GitHub Action:

```
name: Terraform

on:
  push:
    branches:
      - main

jobs:
```

```
terraform:
  runs-on: ubuntu-latest
  steps:
    - name: Checkout code
      uses: actions/checkout@v2
    - name: Setup Terraform
      uses: hashicorp/setup-terraform@v1
    - name: Terraform Init
      run: terraform init
    - name: Terraform Plan
      run: terraform plan
```

## Advanced AWS Services with Terraform:

Deploying S3 buckets and object storage

### Step 1: Create an S3 Bucket

Add the following configuration to create an S3 bucket in main.tf:

```
resource "aws_s3_bucket" "example" {
  bucket = "my-terraform-bucket-12345" # Replace with a unique bucket name
  acl    = "private"

  tags = {
    Name      = "MyS3Bucket"
    Environment = "Dev"
  }
}
```

Initialize Terraform:

```
terraform init
```

Apply the configuration:

```
terraform apply
```

### Step 2: Enable Versioning and Logging

Update the configuration to enable versioning and logging:

```
resource "aws_s3_bucket" "example" {
  bucket = "my-terraform-bucket-12345"
  acl    = "private"

  versioning {
    enabled = true
  }

  logging {
    target_bucket = "my-log-bucket"
    target_prefix = "logs/"
  }
}
```

```
tags = {
  Name      = "MyS3Bucket"
  Environment = "Dev"
}
```

Apply the updated configuration:

```
terraform apply
```

### Step 3: Add an S3 Object

Add an object to the bucket:

```
resource "aws_s3_object" "example" {
  bucket = aws_s3_bucket.example.id
  key    = "example.txt"
  content = "Hello, Terraform!"
}
```

Apply the changes:

```
terraform apply
```

Creating IAM roles and policies

### Step 1: Create an IAM Role

Add the following configuration:

```
resource "aws_iam_role" "example" {
  name = "example-role"

  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect = "Allow"
      Principal = {
        Service = "ec2.amazonaws.com"
      }
      Action = "sts:AssumeRole"
    }]
  })
}
```

Apply the configuration:

```
terraform apply
```

## Step 2: Create an IAM Policy

Add the following configuration:

```
resource "aws_iam_policy" "example_policy" {
  name      = "example-policy"
  description = "A policy to allow S3 access"
  policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Action = "s3:*"
      Effect = "Allow"
      Resource = "*"
    }]
  })
}
```

Apply the configuration:

```
terraform apply
```

## Step 3: Attach the Policy to the Role

Add the following:

```
resource "aws_iam_role_policy_attachment" "example" {
  role      = aws_iam_role.example.name
  policy_arn = aws_iam_policy.example_policy.arn
}
```

Apply the configuration:

```
terraform apply
```

Configuring security groups and network ACLs

## Step 1: Create a Security Group

Add the following configuration:

```
resource "aws_security_group" "example" {
  name      = "example-sg"
  description = "Allow SSH and HTTP"
  vpc_id    = aws_vpc.example.id

  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

```
}

egress {
  from_port = 0
  to_port   = 0
  protocol  = "-1"
  cidr_blocks = ["0.0.0.0/0"]
}

tags = {
  Name = "ExampleSecurityGroup"
}
}
```

Apply the configuration:

```
terraform apply
```

## Step 2: Create a Network ACL

Add the following configuration:

```
resource "aws_network_acl" "example" {
  vpc_id = aws_vpc.example.id

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    rule_no   = 100
    cidr_block = "0.0.0.0/0"
    action    = "allow"
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    rule_no   = 200
    cidr_block = "0.0.0.0/0"
    action    = "allow"
  }

  tags = {
    Name = "ExampleNetworkACL"
  }
}
```

Apply the configuration:

```
terraform apply
```

## Automating deployments with Terraform Cloud

### Step 1: Sign Up for Terraform Cloud

- Create an account at Terraform Cloud.
- Create a new workspace.

### Step 2: Configure Terraform Cloud Backend

Add the following backend configuration in main.tf:

```
terraform {  
  backend "remote" {  
    organization = "your-organization-name"  
  
    workspaces {  
      name = "example-workspace"  
    }  
  }  
}
```

Initialize Terraform:

```
terraform init
```

### Step 3: Push Code to Version Control

Commit your Terraform configuration:

```
git init  
git add .  
git commit -m "Initial commit"  
git push origin main
```

\*Connect the Terraform Cloud workspace to your Git repository.

### Step 4: Automate Plan and Apply

- Configure Terraform Cloud to automatically run plan and apply when changes are pushed to the repository.
- Monitor the deployment in the Terraform Cloud UI.

## Hands-on Lab: Creating an S3 Bucket with Versioning

Write a Terraform Configuration File to Create an S3 Bucket with Versioning

### Step 1: Set Up the Project Directory

Create a new directory for your Terraform project:

```
mkdir terraform-s3-bucket  
cd terraform-s3-bucket
```

Create the main Terraform configuration file:

```
touch main.tf
```



## Step 2: Add the Configuration in main.tf

Here's the Terraform configuration to create an S3 bucket with versioning:

```
# Configure the AWS Provider
provider "aws" {
  region = "us-east-1" # Replace with your preferred region
}

# Create the S3 Bucket
resource "aws_s3_bucket" "example" {
  bucket = "my-versioned-bucket-12345" # Replace with a unique bucket name
  acl    = "private"                  # Set bucket ACL to private

  tags = {
    Name      = "VersionedS3Bucket"
    Environment = "Dev"
  }
}

# Enable Versioning on the Bucket
resource "aws_s3_bucket_versioning" "example" {
  bucket = aws_s3_bucket.example.id

  versioning_configuration {
    status = "Enabled" # Enable versioning
  }
}

# Add Output to Display the Bucket Name
output "bucket_name" {
  value = aws_s3_bucket.example.bucket
}
```

Apply the Configuration to Create the S3 Bucket

### Step 1: Initialize Terraform

Run the following command to initialize the Terraform project and download the necessary provider plugins:

```
terraform init
```

Expected Output:

Terraform has been successfully initialized!

### Step 2: Validate the Configuration

Check if the configuration is valid:

```
terraform validate
```

Expected Output:

Success! The configuration is valid.

### Step 3: Plan the Infrastructure Changes

Preview the changes Terraform will make:

```
terraform plan
```

Expected Output:

A detailed summary of resources to be created, including the S3 bucket and its versioning configuration.

### Step 4: Apply the Configuration

Run the following command to create the S3 bucket:

```
terraform apply
```

\*When prompted, type yes to confirm.

Expected Output:

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

### Verify the S3 Bucket in AWS

- a) Log in to the AWS Management Console.
- b) Navigate to the S3 service.
- c) Confirm that:
  - ✓ The bucket (my-versioned-bucket-12345) has been created.
  - ✓ Versioning is enabled under the Properties tab for the bucket.

### Clean Up Resources (Optional)

To avoid unnecessary costs, destroy the resources when you're done:

```
terraform destroy
```

Expected Output:

```
Destroy complete! Resources: 2 destroyed.
```

## Terraform and AWS Security:

Implementing security best practices in Terraform configurations

### Best Practice 1: Avoid Hardcoding Secrets

Problem: Storing sensitive data like API keys or passwords directly in configuration files.

Solution: Use environment variables or secret management tools.

Example Using Environment Variables:

```
provider "aws" {  
  region    = "us-east-1"  
  access_key = var.aws_access_key  
  secret_key = var.aws_secret_key  
}  
  
variable "aws_access_key" {}  
variable "aws_secret_key" {}
```

Run Terraform with environment variables:

```
export TF_VAR_aws_access_key="your-access-key"
export TF_VAR_aws_secret_key="your-secret-key"
terraform apply
```

### Best Practice 2: Use Remote Backends to Store State

Problem: Local .tfstate files can be accessed or modified accidentally.

Solution: Use remote backends like S3 with encryption.

Example Using S3 as a Backend:

```
terraform {
  backend "s3" {
    bucket    = "my-terraform-state-bucket"
    key       = "state/terraform.tfstate"
    region    = "us-east-1"
    encrypt   = true
    dynamodb_table = "terraform-locking"
  }
}
```

Initialize Terraform with the backend:

```
terraform init
```

### Best Practice 3: Scan Configurations for Security Issues

Use tools like tflint, tfsec, or checkov to identify vulnerabilities.

```
tflint
tfsec .
```

### Best Practice 4: Use Least Privilege Principle for IAM Roles

Grant minimal permissions required for Terraform and AWS resources (discussed in Section 2).

Securing AWS resources with IAM roles and policies

#### **Step 1: Create an IAM Role**

Configuration:

```
resource "aws_iam_role" "example" {
  name = "example-role"

  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect = "Allow"
      Principal = {
        Service = "ec2.amazonaws.com"
      }
      Action = "sts:AssumeRole"
    }]
  })
}
```

## Step 2: Create a Policy for S3 Access

Configuration:

```
resource "aws_iam_policy" "example_policy" {
  name      = "example-policy"
  description = "Allow S3 access"
  policy = jsonencode({
    Version = "2012-10-17"
    Statement = [{
      Effect = "Allow"
      Action = ["s3:ListBucket", "s3:GetObject", "s3:PutObject"]
      Resource = ["arn:aws:s3:::example-bucket", "arn:aws:s3:::example-bucket/*"]
    }]
  })
}
```

## Step 3: Attach the Policy to the Role

Configuration:

```
resource "aws_iam_role_policy_attachment" "example_attachment" {
  role      = aws_iam_role.example.name
  policy_arn = aws_iam_policy.example_policy.arn
}
```

## Step 4: Verify and Apply the Configuration

Initialize Terraform:

```
terraform init
```

Plan and apply:

```
terraform plan
terraform apply
```

Using security groups to control network traffic

Security groups act as virtual firewalls to control inbound and outbound traffic to AWS resources.

## Step 1: Create a Security Group

Configuration:

```
resource "aws_security_group" "example" {
  name      = "example-sg"
  description = "Allow SSH and HTTP"
  vpc_id    = aws_vpc.example.id

  # Allow inbound SSH traffic
  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

```
# Allow inbound HTTP traffic
ingress {
  from_port = 80
  to_port   = 80
  protocol  = "tcp"
  cidr_blocks = ["0.0.0.0/0"]
}

# Allow all outbound traffic
egress {
  from_port = 0
  to_port   = 0
  protocol  = "-1"
  cidr_blocks = ["0.0.0.0/0"]
}

tags = {
  Name = "ExampleSecurityGroup"
}
```

## Step 2: Associate Security Group with EC2 Instance

Configuration:

```
resource "aws_instance" "example" {
  ami          = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
  security_groups = [
    aws_security_group.example.name
  ]

  tags = {
    Name = "ExampleInstance"
  }
}
```

## Step 3: Verify and Apply the Configuration

Initialize Terraform:

```
terraform init
```

Plan and apply:

```
terraform plan
terraform apply
```

## Real-world Use Cases:

Deploying multi-tier applications on AWS

Objective: Deploy a simple multi-tier application with a web tier (EC2) and a database tier (RDS).

### Step 1: Define the Project Structure

Directory Layout:

```
multi-tier-app/  
├── main.tf      # Main configuration  
├── variables.tf  # Variables for reusability  
└── outputs.tf   # Outputs to access resources
```

### Step 2: Create the VPC and Subnets

Add the following configuration in main.tf:

```
provider "aws" {  
  region = "us-east-1"  
}  
  
resource "aws_vpc" "app_vpc" {  
  cidr_block      = "10.0.0.0/16"  
  enable_dns_support = true  
  enable_dns_hostnames = true  
  tags = {  
    Name = "AppVPC"  
  }  
}  
  
resource "aws_subnet" "public_subnet" {  
  vpc_id      = aws_vpc.app_vpc.id  
  cidr_block   = "10.0.1.0/24"  
  map_public_ip_on_launch = true  
  availability_zone = "us-east-1a"  
  tags = {  
    Name = "PublicSubnet"  
  }  
}  
  
resource "aws_subnet" "private_subnet" {  
  vpc_id      = aws_vpc.app_vpc.id  
  cidr_block   = "10.0.2.0/24"  
  availability_zone = "us-east-1a"  
  tags = {  
    Name = "PrivateSubnet"  
  }  
}
```

### Step 3: Launch the Web Tier

- Add an EC2 instance in the public subnet for the web server.
- Create a security group to allow HTTP and SSH traffic.

```
resource "aws_security_group" "web_sg" {
  vpc_id = aws_vpc.app_vpc.id
  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

resource "aws_instance" "web_server" {
  ami          = "ami-0c55b159cbfafa1f0"
  instance_type = "t2.micro"
  subnet_id    = aws_subnet.public_subnet.id
  security_groups = [aws_security_group.web_sg.name]
  tags = {
    Name = "WebServer"
  }
}
```

### Step 4: Launch the Database Tier

Use RDS in the private subnet for the database tier.

```
resource "aws_db_instance" "db_instance" {
  allocated_storage = 20
  engine            = "mysql"
  engine_version    = "8.0"
  instance_class     = "db.t2.micro"
  name              = "appdb"
  username          = "admin"
  password          = "password123"
  vpc_security_group_ids = [aws_security_group.web_sg.id]
  db_subnet_group_name = aws_db_subnet_group.db_subnet_group.name
}

resource "aws_db_subnet_group" "db_subnet_group" {
```

```
name      = "db-subnet-group"
subnet_ids = [aws_subnet.private_subnet.id]
tags = {
  Name = "DBSubnetGroup"
}
```

### Step 5: Apply the Configuration

Initialize Terraform:

```
terraform init
```

Plan the infrastructure:

```
terraform plan
```

Apply the configuration:

```
terraform apply
```

Automating infrastructure for CI/CD pipelines

Objective: Create a CI/CD pipeline with CodePipeline, CodeBuild, and S3.

#### Step 1: Create an S3 Bucket for CodePipeline

```
resource "aws_s3_bucket" "codepipeline_bucket" {
  bucket = "my-codepipeline-bucket"
  acl    = "private"
  tags = {
    Name = "CodePipelineBucket"
  }
}
```

#### Step 2: Define the CodePipeline

```
resource "aws_codepipeline" "my_pipeline" {
  name      = "my-pipeline"
  role_arn = aws_iam_role.codepipeline_role.arn

  artifact_store {
    type = "S3"
    location = aws_s3_bucket.codepipeline_bucket.bucket
  }

  stage {
    name = "Source"

    action {
      name      = "SourceAction"
      category  = "Source"
      owner     = "AWS"
      provider  = "S3"
      version   = "1"
      output_artifacts = ["SourceOutput"]
    }
  }
}
```



```

    configuration = {
      S3Bucket = aws_s3_bucket.codepipeline_bucket.bucket
      S3ObjectKey = "source.zip"
    }
  }
}

stage {
  name = "Build"

  action {
    name      = "BuildAction"
    category  = "Build"
    owner     = "AWS"
    provider  = "CodeBuild"
    input_artifacts = ["SourceOutput"]
    output_artifacts = ["BuildOutput"]

    configuration = {
      ProjectName = aws_codebuild_project.my_project.name
    }
  }
}
}

```

### Step 3: Define the CodeBuild Project

```

resource "aws_codebuild_project" "my_project" {
  name      = "MyProject"
  description = "My CodeBuild Project"
  build_timeout = 5
  service_role = aws_iam_role.codebuild_role.arn

  source {
    type = "S3"
    location = "${aws_s3_bucket.codepipeline_bucket.bucket}/source.zip"
  }

  artifacts {
    type = "NO_ARTIFACTS"
  }

  environment {
    compute_type      = "BUILD_GENERAL1_SMALL"
    image             = "aws/codebuild/standard:5.0"
    type              = "LINUX_CONTAINER"
    privileged_mode    = true
  }
}

```

## Implementing infrastructure as code for disaster recovery

Objective: Implement automated recovery for critical infrastructure using Terraform.

### Step 1: Define Remote State

Use S3 for state management to track the primary infrastructure.

```
terraform {  
  backend "s3" {  
    bucket    = "primary-state"  
    key       = "terraform.tfstate"  
    region    = "us-east-1"  
  }  
}
```

### Step 2: Define Secondary Resources

Use data sources to fetch the current infrastructure.

Create duplicates in a disaster recovery region.

```
provider "aws" {  
  alias = "primary"  
  region = "us-east-1"  
}  
  
provider "aws" {  
  alias = "dr"  
  region = "us-west-2"  
}  
  
data "aws_instance" "primary_instance" {  
  provider = aws.primary  
  instance_id = "i-1234567890abcdef"  
}  
  
resource "aws_instance" "dr_instance" {  
  provider    = aws.dr  
  ami         = data.aws_instance.primary_instance.ami  
  instance_type = data.aws_instance.primary_instance.instance_type  
  tags = {  
    Name = "DRInstance"  
  }  
}
```

### Step 3: Apply Disaster Recovery Plan

Initialize Terraform with the new region:

```
terraform init
```

Apply the disaster recovery configuration:

```
terraform apply
```

# Hands-on Lab: Creating a Multi-Tier Application

## Objective

Deploy a multi-tier application on AWS consisting of:

- Web Tier: EC2 instances behind a Load Balancer.
- Database Tier: RDS MySQL instance.

## Step 1: Define the Project Structure

Directory Layout:

```
multi-tier-app/  
├── main.tf      # Main configuration file  
├── variables.tf # Variables for reusability  
└── outputs.tf   # Outputs to access resources
```

## Step 2: Write the Terraform Configuration File

### Configure the Provider

Add the AWS provider in main.tf:

```
provider "aws" {  
  region = "us-east-1" # Specify your preferred region  
}
```

### Create a VPC and Subnets

Define a VPC with public and private subnets.

```
resource "aws_vpc" "app_vpc" {  
  cidr_block      = "10.0.0.0/16"  
  enable_dns_support = true  
  enable_dns_hostnames = true  
  tags = {  
    Name = "AppVPC"  
  }  
}  
  
resource "aws_subnet" "public_subnet" {  
  vpc_id      = aws_vpc.app_vpc.id  
  cidr_block   = "10.0.1.0/24"  
  map_public_ip_on_launch = true  
  availability_zone = "us-east-1a"  
  tags = {  
    Name = "PublicSubnet"  
  }  
}  
  
resource "aws_subnet" "private_subnet" {  
  vpc_id      = aws_vpc.app_vpc.id  
  cidr_block   = "10.0.2.0/24"  
  availability_zone = "us-east-1a"  
  tags = {  
    Name = "PrivateSubnet"  
  }  
}
```

```
}  
}
```

### Configure the Security Groups

Define security groups for the load balancer, EC2 instances, and RDS.

```
resource "aws_security_group" "web_sg" {  
  vpc_id = aws_vpc.app_vpc.id  
  description = "Allow HTTP and SSH traffic"  
  
  ingress {  
    from_port = 80  
    to_port   = 80  
    protocol  = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
  
  ingress {  
    from_port = 22  
    to_port   = 22  
    protocol  = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
  
  egress {  
    from_port = 0  
    to_port   = 0  
    protocol  = "-1"  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
  
  tags = {  
    Name = "WebSecurityGroup"  
  }  
}  
  
resource "aws_security_group" "db_sg" {  
  vpc_id = aws_vpc.app_vpc.id  
  description = "Allow MySQL traffic"  
  
  ingress {  
    from_port = 3306  
    to_port   = 3306  
    protocol  = "tcp"  
    security_groups = [aws_security_group.web_sg.id]  
  }  
  
  egress {  
    from_port = 0  
    to_port   = 0  
    protocol  = "-1"  
  }  
}
```

```

    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = {
    Name = "DBSecurityGroup"
  }
}

```

### Deploy the EC2 Instances

Launch EC2 instances in the public subnet.

```

resource "aws_instance" "web_server" {
  ami          = "ami-0c55b159cbfafa1f0" # Replace with a valid AMI ID
  instance_type = "t2.micro"
  subnet_id    = aws_subnet.public_subnet.id
  security_groups = [aws_security_group.web_sg.name]

  tags = {
    Name = "WebServer"
  }
}

```

### Deploy a Load Balancer

Create an Application Load Balancer for the EC2 instances.

```

resource "aws_lb" "app_lb" {
  name          = "app-load-balancer"
  internal      = false
  load_balancer_type = "application"
  security_groups = [aws_security_group.web_sg.id]
  subnets      = [aws_subnet.public_subnet.id]

  tags = {
    Name = "AppLoadBalancer"
  }
}

resource "aws_lb_target_group" "web_tg" {
  name     = "web-target-group"
  port     = 80
  protocol = "HTTP"
  vpc_id   = aws_vpc.app_vpc.id
}

resource "aws_lb_listener" "http_listener" {
  load_balancer_arn = aws_lb.app_lb.arn
  port              = 80
  protocol          = "HTTP"

  default_action {
    type = "forward"
    target_group_arn = aws_lb_target_group.web_tg.arn
  }
}

```

```

}
}

resource "aws_lb_target_group_attachment" "web_target" {
  target_group_arn = aws_lb_target_group.web_tg.arn
  target_id       = aws_instance.web_server.id
  port            = 80
}

```

### Deploy an RDS Database

Create an RDS MySQL instance in the private subnet.

```

resource "aws_db_instance" "app_db" {
  allocated_storage = 20
  engine            = "mysql"
  engine_version    = "8.0"
  instance_class    = "db.t2.micro"
  name              = "appdb"
  username          = "admin"
  password          = "password123"
  vpc_security_group_ids = [aws_security_group.db_sg.id]
  db_subnet_group_name = aws_db_subnet_group.db_subnet_group.name
}

resource "aws_db_subnet_group" "db_subnet_group" {
  name      = "db-subnet-group"
  subnet_ids = [aws_subnet.private_subnet.id]
  tags = {
    Name = "DBSubnetGroup"
  }
}

```

### Add Outputs

Define outputs to display key information.

```

output "load_balancer_dns" {
  value = aws_lb.app_lb.dns_name
}

output "database_endpoint" {
  value = aws_db_instance.app_db.endpoint
}

```

## **Step 3: Apply the Configuration to Deploy the Application**

### Initialize Terraform

Run the following command to initialize Terraform and download required providers:

```
terraform init
```

### Validate the Configuration

Validate the syntax and configuration:

```
terraform validate
```

### Plan the Deployment

Generate a detailed plan to verify resources:

```
terraform plan
```

### Apply the Configuration

Deploy the infrastructure:

```
terraform apply
```

\*When prompted, type yes to confirm.

### **Step 4: Verify the Deployment**

Access the Load Balancer DNS:

- a) Copy the DNS name from the output of terraform apply or from the AWS console.
- b) Open it in a browser to verify the web server is accessible.

Verify the RDS database:

- a) Use the endpoint from the Terraform output.
- b) Connect using a MySQL client with the username and password defined in the configuration.